

# CE 473: Artificial Intelligence

Autumn 2017

## A\* Search

first AI robot search algorithm

Dieter Fox

Based on slides from Pieter Abbeel & Dan Klein  
Multiple slides from Stuart Russell, Andrew Moore, Luke Zettlemoyer


## Today

- A\* Search
- Heuristic Design
- Graph search



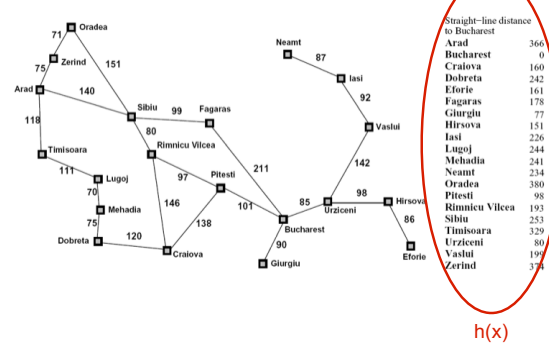
## What is a *Heuristic*?

- An *estimate* of how close a state is to a goal
- Designed for a particular search problem



- Examples: Manhattan distance:  $10+5 = 15$   
Euclidean distance: 11.2

## Example: Heuristic Function





## Greedy Search



## Best First (Greedy)

- Strategy: expand a node that you think is closest to a goal state
  - Heuristic: estimate of distance to nearest goal for each state
- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS






## Optimality of A\* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will exit the fringe before B




## Optimality of A\* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B

1.  $f(n)$  is less or equal to  $f(A)$

$$\begin{aligned} f(n) &= g(n) + h(n) && \text{Definition of f-cost} \\ f(n) &\leq g(A) && \text{Admissibility of h} \\ g(A) &= f(A) && h = 0 \text{ at a goal} \end{aligned}$$




## Optimality of A\* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor  $n$  of A is on the fringe, too (maybe A!)
- Claim:  $n$  will be expanded before B

1.  $f(n)$  is less or equal to  $f(A)$
2.  $f(A)$  is less than  $f(B)$

$$\begin{aligned} g(A) &< g(B) && \text{B is suboptimal} \\ f(A) &< f(B) && h = 0 \text{ at a goal} \end{aligned}$$



## Optimality of A\* Tree Search


Proof:

- Imagine B is on the fringe
  - Some ancestor  $n$  of A is on the fringe, too (maybe A!)
  - Claim:  $n$  will be expanded before B
1.  $f(n)$  is less or equal to  $f(A)$
  2.  $f(A)$  is less than  $f(B)$
  3.  $n$  expands before B

- All ancestors of A expand before B  
[since  $n$  will put the next node  $n'$  toward A on the fringe; repeat claim for  $n'$  until you hit A]


- A expands before B
- A\* search is optimal

$$f(n) \leq f(A) < f(B)$$




## UCS vs A\* Contours

- Uniform-cost expanded in all directions
- A\* expands mainly toward the goal, but hedges its bets to ensure optimality




## Which Algorithm?

- Uniform cost search (UCS):




## Which Algorithm?

- A\*, Manhattan Heuristic:




## Which Algorithm?

- Best First / Greedy, Manhattan Heuristic:



## Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available



- Inadmissible heuristics are often useful too

## Creating Heuristics

8-puzzle:

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- What are the states?
- How many states? **9!**
- What are the actions? **2, 3, 4**
- What states can I reach from the start state?
- What should the costs be?

## 8 Puzzle I

- Heuristic: Number of tiles misplaced
- $h(\text{start}) = 8$
- Is it admissible?

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

	Average nodes expanded when optimal path has length...		
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	$3.6 \times 10^6$
TILES	13	39	227

## 8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total Manhattan distance
- $h(\text{start}) = 3 + 1 + 2 + \dots$

= 18

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- Admissible?

	Average nodes expanded when optimal path has length...		
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73


make it more similar to actually cost

## 8 Puzzle III

- How about using the *actual cost* as a heuristic?
  - Would it be admissible? ✓
  - Would we save on nodes expanded?
- What's wrong with it?
- With A\*: a trade-off between quality of estimate and work per node!

## Trivial Heuristics, Dominance

- Dominance:  $\underline{h}_a \geq h_c$  if  $\forall n : h_a(n) \geq h_c(n)$
- Heuristics form a semi-lattice:
  - Max of admissible heuristics is admissible
- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic




## A\* Applications

- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- ...


## Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work. Why?



## Graph Search


- In BFS, for example, we shouldn't bother expanding some nodes (which, and why?)




## Graph Search

- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states ("closed set")
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Hint: in python, store the closed set as a **set**, not a list
- Can graph search wreck completeness? Why/why not? ✓
- How about optimality?

## A\* Graph Search Gone Wrong



## Consistency of Heuristics



## Optimality of A\* Graph Search


- Sketch: consider what A\* does with a consistent heuristic:

- Nodes are popped with non-decreasing  $f$ -scores: for all  $n, n'$  with  $n'$  popped after  $n$ :  $f(n') \geq f(n)$

- Proof by induction: (1) always pop the lowest  $f$ -score from the fringe, (2) all new nodes have larger (or equal) scores, (3) add them to the fringe, (4) repeat!

- For every state  $s$ , nodes that reach  $s$  optimally are expanded before nodes that reach  $s$  sub-optimally

- Result: A\* graph search is optimal by consistency



## Optimality

- Tree search:
    - A\* optimal if heuristic is admissible (and non-negative)
    - UCS is a special case ( $h = 0$ )
  - Graph search:
    - A\* optimal if heuristic is consistent
    - UCS optimal ( $h = 0$  is consistent)
  - Consistency implies admissibility
  - In general, natural admissible heuristics tend to be consistent, especially if from relaxed problems
- typically are same (generally could be not)

## Summary: A\*

- A\* uses both backward costs and (estimates of) forward costs
- A\* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems